

Yash Srivastava

Phone: +91-9267969254 | Email: ys7.work@gmail.com | GitHub | LinkedIn

Professional Summary

Applied AI Engineer with 2+ years of experience building and productionizing Agentic AI, multi-agent, MCP, RAG/GraphRAG, and LLM-powered systems. Strong backend and platform engineering foundation across APIs, distributed services, asynchronous workflows, data platforms, cloud infrastructure, CI/CD, and end-to-end production ownership.

Technical Skills

Languages: Python, Java, TypeScript, JavaScript, SQL, YAML

Generative AI & LLMs: GenAI, LLM Integration, Prompt Engineering, Fine-tuning (LoRA/SFT), RAG, GraphRAG, Agentic AI, Multi-Agent Systems, MCP Servers, LangChain, LangGraph, CrewAI, Spring AI, Graphiti, Google Gemini, OpenAI, Claude, Azure OpenAI, Ollama, LiteLLM, Langfuse

Backend & Integrations: FastAPI, Django, SQLAlchemy, Alembic, Celery, REST APIs, Microservices, Playwright, Spring Framework, Spring Boot, Slack Socket Mode, Meta Graph API, Pydantic, OAuth 2.0, JWT, RBAC

Databases & BI: PostgreSQL (pgvector), Neo4j, MySQL/Aurora, MongoDB/DocumentDB, Redis, Elasticsearch, Prometheus, Azure SQL Database, Apache Superset, AWS QuickSight, Azure Data Factory, BM25, HNSW, Hybrid Search, Reranking

Infrastructure & DevOps: Linux, Docker, Docker Compose, Azure, AWS (EC2, S3, ECS Fargate, RDS, Secrets Manager, IAM), CI/CD, GitHub Actions, Jenkins, LocalStack, Valkey, RunPod, Modal, Grafana, Kafka, Nginx, Kubernetes

Automation: n8n

Professional Experience

1.) ArmorCode Inc.

Aug 2025 – Present

Applied AI Engineer

Gurugram, India

- Delivered 30+ production AI automations and 12+ MCP servers across Engineering, Security, Support, PM, IT, Sales, DevOps, QA, Documentation, and RevOps functions, enabling reusable AI agents and AI-driven workflows, eliminating 100+ manual hours
- Grounded enterprise AI agents in 1M+ product and organizational knowledge entities by architecting the central GraphRAG knowledge layer with LLM entity and relationship extraction, Neo4j traversal, pgvector HNSW search, hybrid retrieval, Redis graph caching, and LiteLLM orchestration.
- Designed Office of Agents, Office of Agents enabling 10+ business and engineering function to securely interact with organization-specific AI orchestrators by productionizing the platform, supporting autonomous, human-in-the-loop agentic applications executions and orchestrating through hierarchical delegation, persistent state, scheduled execution and governance controls.
- Reduced AI inference overhead and coding-agent complexity for office of agents, by integrating, Headroom for context compression and Ponytail for coding-agent optimization in production, steering agents toward simpler implementations, fewer dependencies, native capabilities, and more efficient tool usage.
- Unified business analytics across 4 production AWS accounts and 10+ SaaS sources by replacing AWS QuickSight with a self-hosted PostgreSQL and Apache Superset data platform via cross-account S3 replication, exposing business-data insights, with exposure across superset MCP.
- Enabled repository-level AI engineering across 8 production repositories by extending OpenCode's architecture, with 8+ specialized AI agents and 7+ reusable skills for codebase search, documentation, feature-flag cataloging, and tenant-configuration analysis.
- Owned entire production AI automation infrastructure by managing and modernizing the n8n platform through containerized deployments, multi-stage CI/CD pipelines, Lua-based RBAC, bidirectional GitHub synchronization, controlled promotion, and production-safe upgrades.

2.) Xansr Software Pvt. Ltd.

Feb 2024 – Aug 2025

Junior GenAI Specialist

Remote

- Achieved 90% chatbot accuracy with 100% test coverage by deploying a production GenAI SQL-RAG chatbot using FastAPI, LangChain, and Azure OpenAI, while supporting real-time data ingestion and contextual conversations.
- Streamlined end-to-end production data processing by implementing ETL pipelines on Azure SQL Database, designing SQLAlchemy-based schemas, and integrating Azure Logic Apps for automated data ingestion workflows.
- Improved domain-specific model performance through 2 fine-tuning approaches—supervised fine-tuning and LoRA—on serverless GPU infrastructure, optimizing inference latency, accuracy, and cost efficiency.
- Extended AI applications to voice-based interaction by integrating ElevenLabs text-to-speech capabilities, enabling natural-language voice responses from AI-generated content.

Projects

Creator OS — Multi-Tenant Full-Stack AI Platform

2026

React, TypeScript, Vite, Tailwind CSS, Supabase, FastAPI, PostgreSQL, async SQLAlchemy, Celery, Valkey, Docker, Meta APIs, Hugging Face, Self-Hosted Vaidya Model

- Developed a multi-tenant Creator OS supporting 9+ platform capabilities, including Instagram OAuth, Meta Graph API integrations, AI creator intelligence, content planning, brand matchmaking, outreach, email automation, media storage, and AI fallbacks using FastAPI, PostgreSQL, Celery/Valkey, and Docker.

OpenClaw NSE Market Intelligence Agent

2025

Python, OpenClaw Agent Framework, Google Gemini, Groww API, Yahoo Finance, Finnhub, Telegram Bot API

- Engineered an autonomous NSE market intelligence agent integrating 6 external AI, market-data, and communication services to provide real-time market analysis, sentiment scoring, 24/7 portfolio monitoring, NL recommendations, and human-in-the-loop trade execution.

Education

The NorthCap University

Aug 2020 – Jun 2024

- B.Tech, Computer Science Engineering

- Specialization: AI & ML